Solution 1: Bayesian Linear Model

The posterior distribution is obtained by Bayes' rule

$$\underbrace{p(\boldsymbol{\theta}|\mathbf{X},\mathbf{y})}_{\text{posterior}} = \frac{\overbrace{p(\mathbf{y}|\mathbf{X},\boldsymbol{\theta})}^{\text{likelihood}}\overbrace{q(\boldsymbol{\theta})}^{\text{prior}}}{\underbrace{p(\mathbf{y}|\mathbf{X})}_{\text{marginal}}}.$$

In the Bayesian linear model we have a Gaussian likelihood: $\mathbf{y} \mid \mathbf{X}, \boldsymbol{\theta} \sim \mathcal{N}(\mathbf{X}\boldsymbol{\theta}, \sigma^2 \boldsymbol{I}_n)$, i.e.,

$$\begin{aligned} p(\mathbf{y}|\mathbf{X},\boldsymbol{\theta}) &\propto \exp\left[-\frac{1}{2\sigma^2}(\mathbf{y}-\mathbf{X}\boldsymbol{\theta})^\top(\mathbf{y}-\mathbf{X}\boldsymbol{\theta})\right] \\ &= \exp\left[-\frac{\|\mathbf{y}-\mathbf{X}\boldsymbol{\theta}\|_2^2}{2\sigma^2}\right] \\ &= \exp\left[-\frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top\mathbf{x}^{(i)})^2}{2\sigma^2}\right]. \end{aligned}$$

Moreover, note that the maximum a posteriori estimate of $\boldsymbol{\theta}$, which is defined by

$$\hat{\boldsymbol{\theta}} = \arg\max_{\boldsymbol{\theta}} p(\boldsymbol{\theta}|\mathbf{X},\mathbf{y})$$

can also be defined by

$$\hat{\boldsymbol{\theta}} = \arg\max_{\boldsymbol{\theta}} \log\left(p(\boldsymbol{\theta}|\mathbf{X},\mathbf{y})\right),$$

since log is a monotonically increasing function, so the maximizer is the same.

(a) If the prior distribution is a uniform distribution over the parameter vectors $\boldsymbol{\theta}$, i.e.,

$$q(\boldsymbol{\theta}) \propto 1,$$

then

$$\begin{aligned} p(\boldsymbol{\theta}|\mathbf{X},\mathbf{y}) &\propto p(\mathbf{y}|\mathbf{X},\boldsymbol{\theta})q(\boldsymbol{\theta}) \\ &\propto \exp\left[-\frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top\mathbf{x}^{(i)})^2}{2\sigma^2}\right]. \end{aligned}$$

With this,

$$\begin{aligned} \hat{\boldsymbol{\theta}} &= \arg\max_{\boldsymbol{\theta}} \log\left(p(\boldsymbol{\theta}|\mathbf{X},\mathbf{y})\right) \\ &= \arg\max_{\boldsymbol{\theta}} -\frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top\mathbf{x}^{(i)})^2}{2\sigma^2} \\ &= \arg\min_{\boldsymbol{\theta}} \frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top\mathbf{x}^{(i)})^2}{2\sigma^2} \\ &= \arg\min_{\boldsymbol{\theta}} \sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top\mathbf{x}^{(i)})^2, \qquad (2\sigma^2 \text{ is just a constant scaling}) \end{aligned}$$

so the maximum a posteriori estimate coincides with the empirical risk minimizer for the L2-loss (over the linear models).

(b) If we choose a Gaussian distribution over the parameter vectors $\boldsymbol{\theta}$ as the prior belief, i.e.,

$$q(\boldsymbol{\theta}) \propto \exp\left[-\frac{1}{2\tau^2}\boldsymbol{\theta}^\top\boldsymbol{\theta}\right], \qquad \tau > 0,$$

then

$$\begin{aligned} p(\boldsymbol{\theta}|\mathbf{X},\mathbf{y}) &\propto p(\mathbf{y}|\mathbf{X},\boldsymbol{\theta})q(\boldsymbol{\theta}) \\ &\propto \exp\left[-\frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top \mathbf{x}^{(i)})^2}{2\sigma^2} - \frac{1}{2\tau^2}\boldsymbol{\theta}^\top\boldsymbol{\theta}\right] \\ &= \exp\left[-\frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top \mathbf{x}^{(i)})^2}{2\sigma^2} - \frac{\|\boldsymbol{\theta}\|_2^2}{2\tau^2}\right] \end{aligned}$$

With this,

$$\begin{aligned} \hat{\boldsymbol{\theta}} &= \arg\max_{\boldsymbol{\theta}} \log\left(p(\boldsymbol{\theta}|\mathbf{X},\mathbf{y})\right) \\ &= \arg\max_{\boldsymbol{\theta}} -\frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top \mathbf{x}^{(i)})^2}{2\sigma^2} - \frac{\|\boldsymbol{\theta}\|_2^2}{2\tau^2} \\ &= \arg\min_{\boldsymbol{\theta}} \frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top \mathbf{x}^{(i)})^2}{2\sigma^2} + \frac{\|\boldsymbol{\theta}\|_2^2}{2\tau^2} \\ &= \arg\min_{\boldsymbol{\theta}} \sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top \mathbf{x}^{(i)})^2 + \frac{\sigma^2}{\tau^2}\|\boldsymbol{\theta}\|_2^2, \end{aligned}$$

so the maximum a posteriori estimate coincides for the choice of $\lambda = \frac{\sigma^2}{\tau^2} > 0$ with the regularized empirical risk minimizer for the L2-loss with L2 penalty (over the linear models), i.e., the Ridge regression.

(c) If we choose a Laplace distribution over the parameter vectors $\boldsymbol{\theta}$ as the prior belief, i.e.,

$$q(\boldsymbol{\theta}) \propto \exp\left[-\frac{\sum_{i=1}^p |\boldsymbol{\theta}_i|}{\tau}\right], \qquad \tau > 0,$$

then

$$\begin{aligned} p(\boldsymbol{\theta}|\mathbf{X},\mathbf{y}) &\propto p(\mathbf{y}|\mathbf{X},\boldsymbol{\theta})q(\boldsymbol{\theta}) \\ &\propto \exp\left[-\frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top \mathbf{x}^{(i)})^2}{2\sigma^2} - \frac{\sum_{i=1}^p |\boldsymbol{\theta}_i|}{\tau}\right] \\ &= \exp\left[-\frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top \mathbf{x}^{(i)})^2}{2\sigma^2} - \frac{\|\boldsymbol{\theta}\|_1}{\tau}\right] \end{aligned}$$

With this,

$$\begin{aligned} \hat{\boldsymbol{\theta}} &= \arg\max_{\boldsymbol{\theta}} \log\left(p(\boldsymbol{\theta}|\mathbf{X},\mathbf{y})\right) \\ &= \arg\max_{\boldsymbol{\theta}} -\frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top \mathbf{x}^{(i)})^2}{2\sigma^2} - \frac{\|\boldsymbol{\theta}\|_1}{\tau} \\ &= \arg\min_{\boldsymbol{\theta}} \frac{\sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top \mathbf{x}^{(i)})^2}{2\sigma^2} + \frac{\|\boldsymbol{\theta}\|_1}{\tau} \\ &= \arg\min_{\boldsymbol{\theta}} \sum_{i=1}^n (y^{(i)}-\boldsymbol{\theta}^\top \mathbf{x}^{(i)})^2 + \frac{2\sigma^2}{\tau}\|\boldsymbol{\theta}\|_1, \end{aligned}$$

so the maximum a posteriori estimate coincides for the specific choice of $\lambda = \frac{2\sigma^2}{\tau}$ with the regularized empirical risk minimizer for the L2-loss with L1 penalty (over the linear models), i.e., the Lasso regression.

Solution 2: Gaussian Posterior Process

(a) Prior distribution (assuming the same notation as in the lecture):

$$\boldsymbol{f} \sim \mathcal{N}(\boldsymbol{m}, \boldsymbol{K})$$

with $\boldsymbol{m} = m(\boldsymbol{x})$ and $\boldsymbol{K}$ defined by the entries $\boldsymbol{K}_{ij} = k(x_i, x_j)$. NB: Note the (in-)finite Gaussian property of a GP.

(b) Note that the posterior distribution $\boldsymbol{f}|\boldsymbol{y},\boldsymbol{x}$ in this case is different from the one of $\boldsymbol{f}_*|\boldsymbol{x}_*,\boldsymbol{x},\boldsymbol{y}$ and also from the marginal distribution of $\boldsymbol{y} \sim \mathcal{N}(\boldsymbol{m}, \boldsymbol{K} + \sigma^2\boldsymbol{I})$! We have:

$$\begin{aligned} p(\boldsymbol{f}|\boldsymbol{y}) &\propto p(\boldsymbol{y}|\boldsymbol{f}) \cdot p(\boldsymbol{f}) \\ &\propto \exp(-\frac{1}{2}(\boldsymbol{y}-\boldsymbol{f})^\top(\sigma^2\boldsymbol{I})^{-1}(\boldsymbol{y}-\boldsymbol{f})) \cdot \exp(-\frac{1}{2}(\boldsymbol{f}-\boldsymbol{m})^\top\boldsymbol{K}^{-1}(\boldsymbol{f}-\boldsymbol{m})) \\ &\propto \exp(-\frac{1}{2}\{\boldsymbol{f}^\top \underbrace{((\sigma^2\boldsymbol{I})^{-1} + \boldsymbol{K}^{-1})}_{=:\boldsymbol{K}_{post}^{-1}}\boldsymbol{f} - 2\boldsymbol{f}^\top \underbrace{((\sigma^2\boldsymbol{I})^{-1}\boldsymbol{y} + \boldsymbol{K}^{-1}\boldsymbol{m})}_{=:\tilde{\boldsymbol{f}}}\}) \\ &\propto \exp(-\frac{1}{2}\{\boldsymbol{f}^\top\boldsymbol{K}_{post}^{-1}\boldsymbol{f} - 2\boldsymbol{f}^\top\tilde{\boldsymbol{f}}\}) \end{aligned} \tag{1}$$

by removing all constant factors that do not depend on $\boldsymbol{f}$ as we only need to know the density up to a constant of proportionality. By extending the proportionality, we can get a quadratic form in $\boldsymbol{f}$:

$$\begin{aligned} p(\boldsymbol{f}|\boldsymbol{y}) &\propto \exp(-\frac{1}{2}\{\boldsymbol{f}^\top\boldsymbol{K}_{post}^{-1}\boldsymbol{f} - 2\boldsymbol{f}^\top\tilde{\boldsymbol{f}}\}) \\ &\propto \exp(-\frac{1}{2}\{\boldsymbol{f}^\top\boldsymbol{K}_{post}^{-1}\boldsymbol{f} - 2\boldsymbol{f}^\top\boldsymbol{K}_{post}^{-1}\underbrace{\boldsymbol{K}_{post}\tilde{\boldsymbol{f}}}_{:=\boldsymbol{f}_{post}}\}) \\ &\propto \exp(-\frac{1}{2}(\boldsymbol{f}-\boldsymbol{f}_{post})^\top\boldsymbol{K}_{post}^{-1}(\boldsymbol{f}-\boldsymbol{f}_{post})) \end{aligned} \tag{2}$$

which is the so-called *kernel* of a multivariate normal distribution $\mathcal{N}(\boldsymbol{f}_{post}, \boldsymbol{K}_{post})$, i.e., $\boldsymbol{f}|\boldsymbol{y} \sim \mathcal{N}(\boldsymbol{f}_{post}, \boldsymbol{K}_{post})$.

(c) In order to get the posterior predictive distribution for a new sample x_* from the same data-generating process, we could derive

$$p(y_*|x_*,\boldsymbol{y},\boldsymbol{x}) = \int p(y_*|x_*,\boldsymbol{x},\boldsymbol{y},\boldsymbol{f}) \cdot p(\boldsymbol{f}|\boldsymbol{y},\boldsymbol{x})\, d\boldsymbol{f}.$$

This is feasible but cumbersome. Alternatively, we can make use of the fact that the joint distribution of $\boldsymbol{y}$ and y_* is known (cf. slides on noisy GP):

$$\begin{pmatrix}\boldsymbol{y}\\ y_*\end{pmatrix} \sim \mathcal{N}\left(\begin{pmatrix}\boldsymbol{m}\\ m_*\end{pmatrix}, \begin{pmatrix}\boldsymbol{K}+\sigma^2\boldsymbol{I} & \boldsymbol{K}_*\\ \boldsymbol{K}_*^\top & K_{**}\end{pmatrix}\right),$$

with $m_* = m(x_*)$, $\boldsymbol{K}_* = k(x_*,\boldsymbol{x})$ and $K_{**} = k(x_*,x_*)$. The conditional distribution can then be derived using the rule of conditioning for Gaussian distributions:

$$y_*|x_*,\boldsymbol{x},\boldsymbol{y} \sim \mathcal{N}(m_* + \boldsymbol{K}_*^\top(\boldsymbol{K}+\sigma^2\boldsymbol{I})^{-1}(\boldsymbol{y}-\boldsymbol{m}), K_{**} - \boldsymbol{K}_*^\top(\boldsymbol{K}+\sigma^2\boldsymbol{I})^{-1}\boldsymbol{K}_*).$$

(d) To implement a GP with squared exponential kernel and $\ell = 1$, we need the inverse of $\boldsymbol{K}$. $\boldsymbol{x}$ being a vector implies that we have only one feature and thus the entries of our matrix $\boldsymbol{K}$ are

$$\boldsymbol{K} = \begin{pmatrix}1 & \exp(-0.5(x^{(1)}-x^{(2)})^2)\\ \exp(-0.5(x^{(2)}-x^{(1)})^2) & 1\end{pmatrix}.$$

The inverse of $\boldsymbol{K}$ is then given by

$$\frac{1}{1-\exp(-(x^{(1)}-x^{(2)})^2)}\begin{pmatrix}1 & -\exp(-0.5(x^{(1)}-x^{(2)})^2)\\ -\exp(-0.5(x^{(2)}-x^{(1)})^2) & 1\end{pmatrix}.$$

If we have a noisy GP, we would have to add $\sigma^2\boldsymbol{I}_2$ to $\boldsymbol{K}$ with resulting inverse

$$\boldsymbol{K}_y^{-1} = \frac{1}{(1+\sigma^2)^2-\exp(-(x^{(1)}-x^{(2)})^2)}\begin{pmatrix}1+\sigma^2 & -\exp(-0.5(x^{(1)}-x^{(2)})^2)\\ -\exp(-0.5(x^{(2)}-x^{(1)})^2) & 1+\sigma^2\end{pmatrix}.$$

Assuming a zero mean GP, we can derive $\frac{\partial\boldsymbol{K}_y}{\partial\theta}$ with $\theta = \sigma^2$, which gives us the identity matrix. We can thus maximize the marginal likelihood (slide on *Gaussian Process Training*), by finding σ^2 that yields

$$\operatorname{tr}\left(\boldsymbol{K}_y^{-1}\boldsymbol{y}\boldsymbol{y}^\top\boldsymbol{K}_y^{-1} - \boldsymbol{K}_y^{-1}\right) = 0.$$

This can be solved analytically (though quite tedious). We will use a root-finding function for this. For the posterior predictive distribution we can make use of the results from the previous exercise.

```r
library(kernlab)

# set seed, define n, true (unknown) sigma
set.seed(4212)
n <- 2
sigma <- 1

# define kernel with l = 1
kernel_fun <- function(x)
  kernelMatrix(kernel = rbfdot(sigma = 1/2),
               x = x)
kernel_fun_pred <- function(x,y)
  kernelMatrix(kernel = rbfdot(sigma = 1/2),
               x = x, y = y)

# draw data according to the generating process:
x <- rnorm(n)
K <- kernel_fun(x)
K_y <- K + diag(rep(sigma^2,2))
(y <- t(mvtnorm::rmvnorm(1, sigma = K_y)))

##          [,1]
## [1,] 2.012317
## [2,] 1.866819

# function to find the best sigma^2
root_fun <- function(sigmaSq){
  K_y_inv <- solve(K + diag(rep(sigmaSq,2)))
  0.5*sum(diag(K_y_inv%*%y%*%t(y)%*%K_y_inv - K_y_inv))
}

# get the best sigma
(bestSigmaSq <- uniroot(f = root_fun, interval = c(0,20)))$root

## [1] 1.943684

# plot the optimization problem and best sigma
possible_sigvals <- seq(0.001,20,l=1000)
plot(possible_sigvals, sapply(possible_sigvals, root_fun),
     xlab = expression(sigma^2), ylab = "marginal likelihood derivative",
     pch = 20)
abline(h=0, lty=2)
abline(v=bestSigmaSq$root, lty=2)
```

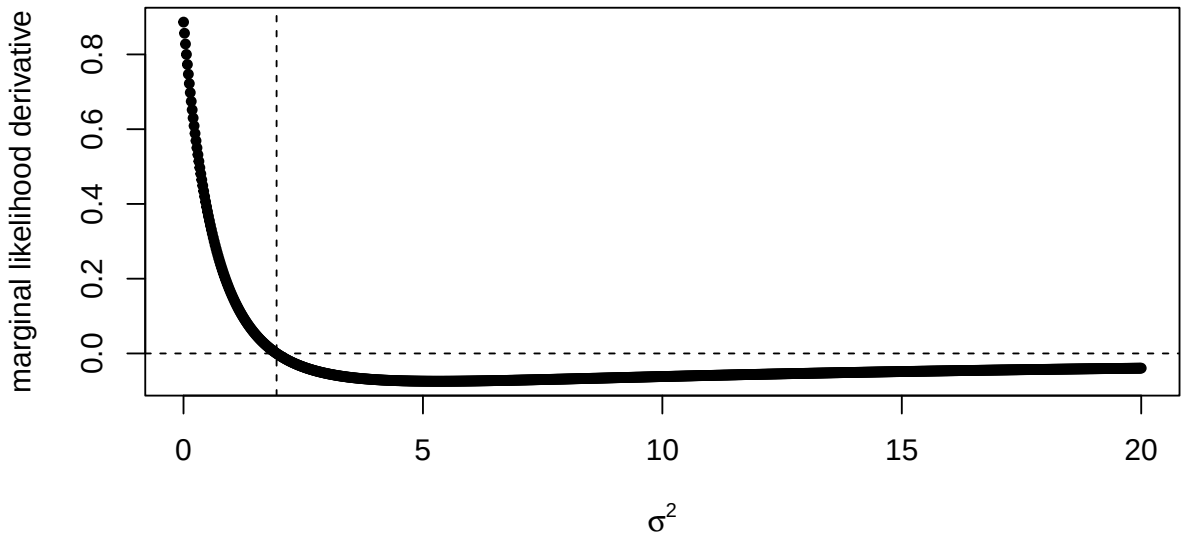


```r
# function to draw samples from the predictive posterior
draw_from_pred_posterior <- function(number_samples, y, x, xstar, sigmaSq = 1)
{

  # invert noisy K
  K_y_inv <- solve(kernel_fun(x) + diag(rep(sigmaSq,2)))
  # get the other K's for new data
  Kstar <- kernel_fun_pred(x,xstar)
  Kstarstar <- kernel_fun(xstar)
  # draw samples according to Ex. (d)
  rnorm(number_samples,
        mean = as.numeric(t(Kstar) %*% K_y_inv %*% y),
        sd = sqrt(as.numeric(Kstarstar - t(Kstar) %*% K_y_inv %*% Kstar))
  )

}

# draw enough samples to get a feeling for the distribution
samples_posterior <-
       draw_from_pred_posterior(number_samples = 1000, sigmaSq = bestSigmaSq$root,
                                y = y, x = x, xstar = 0)
# plot the distribution
hist(samples_posterior, breaks=50, xlab=expression(y["*"]^b))
```

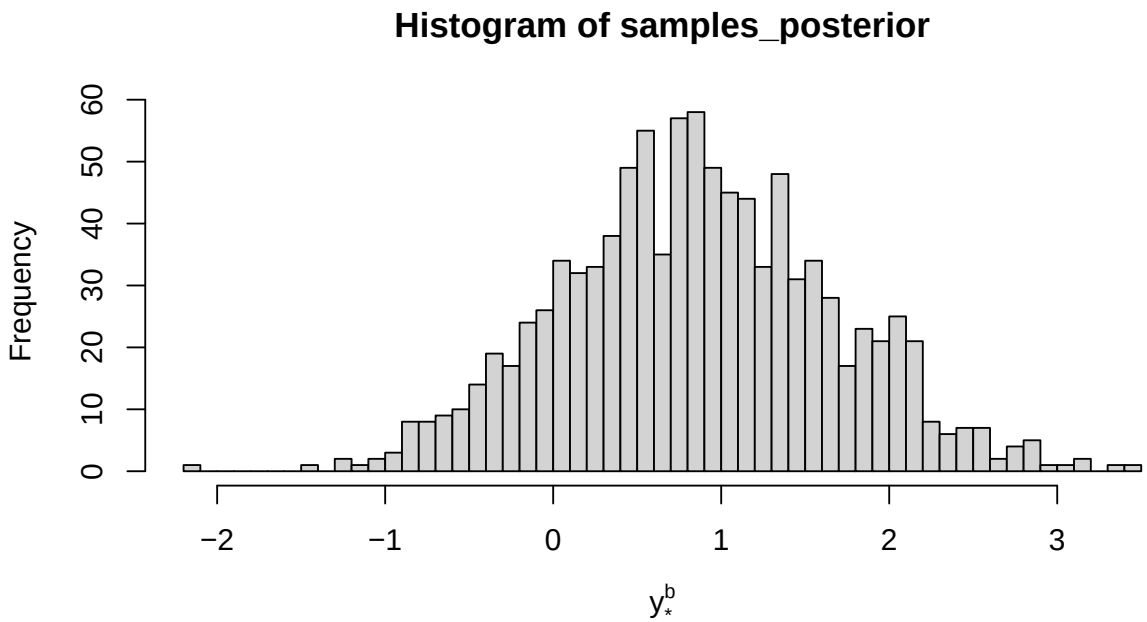
Histogram of samples_posterior
Frequency
0
10
20
30
40
50
60
−2
−1
0
1
2
3
y_*^b