

Solution 1: Cost Curves

- (a) We can simply retrieve the FPRs and TPRs from the table (or matrix) and plot the curves.

```
import numpy as np
import matplotlib.pyplot as plt

# The first column is FPR, the second column is TPR
FPR_TPR_1 = np.array(
    [[0.0, 0.00],
     [0.1, 0.60],
     [0.2, 0.75],
     [0.3, 0.825],
     [0.4, 0.85],
     [0.5, 0.875],
     [0.6, 0.90],
     [0.7, 0.925],
     [0.8, 0.950],
     [0.9, 0.975],
     [1.0, 1.0]]

FPR_TPR_2 = np.array(
    [[0.0, 0.00],
     [0.1, 0.2],
     [0.2, 0.4],
     [0.3, 0.6],
     [0.4, 0.8],
     [0.5, 0.925],
     [0.6, 0.96],
     [0.7, 0.98],
     [0.8, 0.99],
     [0.9, 0.995],
     [1.0, 1.00]]
)

def draw_roc_curves(fpr_tpr_1: np.ndarray, fpr_tpr_2: np.ndarray) -> None:
    fig, ax = plt.subplots(1, 1, figsize=(8, 7))
    ax.plot(fpr_tpr_1[:, 0], fpr_tpr_1[:, 1], marker='o', label='Classifier 1')
    ax.plot(fpr_tpr_2[:, 0], fpr_tpr_2[:, 1], marker='o', label='Classifier 2')

    ax.set_xlabel('False Positive Rate')
    ax.set_ylabel('True Positive Rate')
    ax.grid('on')

    plt.legend()
    plt.savefig('roc_curves.pdf', bbox_inches='tight')
```

The ROC curves are shown in Figure 1

- (b) The cost is computed as

$$\rho_{MCE} = (FNR - FPR) \cdot \pi_+ + FPR$$

For each (FPR, TPR) point in the ROC space, we draw a line of (π_+, ρ_{MCE}) .




Figure 1: ROC Curves.

```
def draw_cost_curves(fpr_tpr_1: np.ndarray, fpr_tpr_2: np.ndarray) -> None:
    fnr_1 = 1 - fpr_tpr_1[:, 1]
    fnr_2 = 1 - fpr_tpr_2[:, 1]
    # Probability of positive
    # Assume P = 50 points in [0.0, 1.0]
    pos_probs = np.linspace(0.0, 1.0, 50)

    # Compute the coefficient (FNR - FPR)
    # Shape after expansion: (N, 1), where N is the number of points in ROC space.
    fnr_minus_fpr_1 = np.expand_dims(fnr_1 - fpr_tpr_1[:, 0], 1)
    fnr_minus_fpr_2 = np.expand_dims(fnr_2 - fpr_tpr_2[:, 0], 1)

    # Costs shape: (N, P)
    costs_1 = fnr_minus_fpr_1 * np.expand_dims(pos_probs, 0) + np.expand_dims(
        fpr_tpr_1[:, 0], 1)
    costs_2 = fnr_minus_fpr_2 * np.expand_dims(pos_probs, 0) + np.expand_dims(
        fpr_tpr_2[:, 0], 1)
    # Point-wise minimum across cost lines.
    cost_curve_1 = costs_1.min(0)
    cost_curve_2 = costs_2.min(0)

    # Find the cross point of two cost curves. The following operation assumes that
    # cost_curve_1 is lower than cost_curve_2 in the beginning of the curves.
    cross_point_ind = np.flatnonzero(cost_curve_1 > cost_curve_2)[0]

    fig, ax = plt.subplots(1, 1, figsize=(8, 7))
    ax.plot(pos_probs, cost_curve_1, label='Classifier 1')
    ax.plot(pos_probs, cost_curve_2, label='Classifier 2')
    ax.vlines(
        pos_probs[cross_point_ind],
        ymin=0.0,
        ymax=cost_curve_1[cross_point_ind],
        linestyle='dashed',
        colors=['k'],
    )
    ax.set_xlabel('Probability of Positive')
    ax.set_ylabel('Error Rate')

    plt.legend()
    plt.savefig('cost_curves.pdf', bbox_inches='tight')
```

```

if __name__ == '__main__':
    draw_roc_curves(FPR_TPR_1, FPR_TPR_2)
    draw_cost_curves(FPR_TPR_1, FPR_TPR_2)

```

Then, we obtain Figure 2.




Figure 2: Cost Curves.

Solution 2: Tomek Links

(a)

```

import matplotlib.pyplot as plt
import numpy as np

from sklearn.datasets import make_classification
from sklearn.metrics import pairwise_distances

def find_tomek_links(x: np.ndarray, y: np.ndarray) -> np.ndarray:
    """Find Tomek Links in samples.

    Args:
        x: Data samples with shape (num_samples, num_features).
        y: Binary class labels with shape (num_samples,).
           1 means positive class, 0 means negative class.

    Returns:
        An array with shape (num_samples,) with binary values,
        for which 1 means that the corresponding sample
        belongs to a Tomek link, while 0 means that the
        sample does not belong to any Tomek link.
    """

    num_samples = x.shape[0]
    dist = pairwise_distances(x)

    # Compute the min. pairwise distance between samples with
    # pos. and neg. samples, respectively. To this end, we employ
    # masked arrays, which can be constructed by
    # applying binary masks to the original array. In the masks, 1

```

```

# means excluding the corresponding samples when performing
# operations e.g. min, argmin, etc.

# pos_mask has shape (num_samples, num_samples).
pos_mask = np.tile(np.expand_dims(1-y, 0), (num_samples, 1))
# Also excluding each sample itself.
# Otherwise, argmin will always point to the sample itself.
pos_mask[np.arange(num_samples), np.arange(num_samples)] = True
dist_pos = np.ma.array(dist, mask=pos_mask)
# Sample indices indicate the closest positive sample to each data sample.
min_dist_inds_pos = dist_pos.argmin(axis=1)

neg_mask = np.tile(np.expand_dims(y, 0), (num_samples, 1))
neg_mask[np.arange(num_samples), np.arange(num_samples)] = True
dist_neg = np.ma.array(dist, mask=neg_mask)
min_dist_inds_neg = dist_neg.argmin(axis=1)

# This interpolation is equivalent to: if y_i = 1, then select the closest pos.
# sample; if y_i = 0, then select the closest neg. sample. In other words, it
# identifies the index of the closest intra-class sample.
min_dist_inds_same_y = min_dist_inds_pos * y + min_dist_inds_neg * (1 - y)
# This interpolation is equivalent to: if y_i = 1, then select the closest neg.
# sample; if y_i = 0, then select the closest pos. sample. In other words, it
# identifies the index of the closest inter-class sample.
min_dist_inds_diff_y = min_dist_inds_pos * (1 - y) + min_dist_inds_neg * y
# Retrieve the distance to the closest intra-class sample and
# inter-class sample
min_dist_same_y = dist[np.arange(num_samples), min_dist_inds_same_y]
min_dist_diff_y = dist[np.arange(num_samples), min_dist_inds_diff_y]

# A sample belongs to a Tomek link if its closest inter-class sample is closer
# than its closest intra-class sample.
tomek_pairs = np.stack([np.arange(num_samples), min_dist_inds_diff_y], axis=1)
# tokek_pairs has shape (num_tomek_pairs, 2) in the end
tomek_pairs = tokek_pairs[min_dist_diff_y <= min_dist_same_y]

# create a binary array and set the elements to 1 if the corresponding indices
# appear in the tokek_pairs.
is_tomek_sample = np.zeros((num_samples, ), dtype=bool)
is_tomek_sample[tomek_pairs.flat] = True
return is_tomek_sample

```

```

(b) def find_kept_samples(
    x: np.ndarray,
    y: np.ndarray,
    is_tomek_sample: np.ndarray
) -> np.ndarray:
    """Given the binary array indicating which sample belongs to a Tomek link. Find
    out which sample should be kept. The sample should be removed if it belongs to
    a Tomek link and it is majority class.

    Args:
        x: Data samples with shape (num_samples, num_features).
        y: Binary class labels with values 0 or 1.
        is_tomek_sample: A binary array with shape (num_samples, ),
            in which 1 means that the sample belongs to a Tomek link.

    Returns:
        A binary array with shape (num_samples, ), in which 1 means that the
        corresponding sample should be kept, otherwise it should be removed.
    """
    num_samples = x.shape[0]
    to_be_kept = np.ones((num_samples, ), dtype=bool)

```

```

# Define which class is majority
num_pos_samples = y.sum()
num_neg_samples = num_samples - num_pos_samples
if num_pos_samples >= num_neg_samples:
    y_major = 1
else:
    y_major = 0

# Only remove the sample in a Tomek link if it is majority class
is_major_tomek_sample = np.logical_and(is_tomek_sample, y == y_major)
to_be_kept[is_major_tomek_sample] = False

return to_be_kept

```

```

(c) def run_experiment() -> None:
    # Generate data
    random_state = np.random.RandomState(11)
    x, y = make_classification(
        100,
        n_features=2,
        n_redundant=0,
        n_clusters_per_class=1,
        flip_y=0.02,
        class_sep=1.0,
        weights=[0.25, 0.75],
        random_state=random_state)

    # Plot the original dataset
    fig, ax = plt.subplots(1, 1)
    x_pos = x[y == 1]
    x_neg = x[y == 0]
    ax.scatter(x_pos[:, 0], x_pos[:, 1], facecolors='none', edgecolors='r', label='
    Pos. samples')
    ax.scatter(x_neg[:, 0], x_neg[:, 1], facecolors='none', edgecolors='b', label='
    Neg. samples')
    ax.set_title('Original Dataset')
    ax.set_xlim(-3, 3)
    ax.set_ylim(-3, 3)
    plt.legend()
    plt.savefig('original_dataset.pdf', bbox_inches='tight')
    plt.clf()

    # Find out which samples belong to Tomek links.
    is_tomek_link = find_tomek_links(x, y)

    fig, ax = plt.subplots(1, 1)
    x_tomek = x[is_tomek_link]
    ax.scatter(x_pos[:, 0], x_pos[:, 1], facecolors='none', edgecolors='r', label='
    Pos. samples')
    ax.scatter(x_neg[:, 0], x_neg[:, 1], facecolors='none', edgecolors='b', label='
    Neg. samples')
    ax.scatter(
        x_tomek[:, 0],
        x_tomek[:, 1],
        facecolors='none',
        edgecolors='g',
        label='Set of Tomek-link samples')
    ax.set_title('Set of Tomek-link samples')
    ax.set_xlim(-3, 3)
    ax.set_ylim(-3, 3)
    plt.legend()
    plt.savefig('tomek_links.pdf', bbox_inches='tight')
    plt.clf()

```

```


# Identify which samples should be kept and which should be removed
to_be_kept = find_kept_samples(x, y, is_tomek_link)

fig, ax = plt.subplots(1, 1)
x_removed = x[~to_be_kept]
x, y = x[to_be_kept], y[to_be_kept]
x_pos, x_neg = x[y == 1], x[y == 0]
ax.scatter(x_pos[:, 0], x_pos[:, 1], facecolors='none', edgecolors='r', label='
Pos. samples')
ax.scatter(x_neg[:, 0], x_neg[:, 1], facecolors='none', edgecolors='b', label='
Neg. samples')
ax.scatter(
    x_removed[:, 0],
    x_removed[:, 1],
    facecolors='none',
    edgecolor='k',
    linestyle='dotted',
    label='Removed samples')
ax.set_title('Majority Samples in Tomek links Removed')
ax.set_xlim(-3, 3)
ax.set_ylim(-3, 3)
plt.legend()
plt.savefig('subsamped_dataset.pdf', bbox_inches='tight')
plt.clf()


if __name__ == '__main__':
    run_experiment()

```


After running the script, we obtain the following figures:



(a) Original Dataset.



(b) Tomek Links.



(c) Subsampled Dataset.