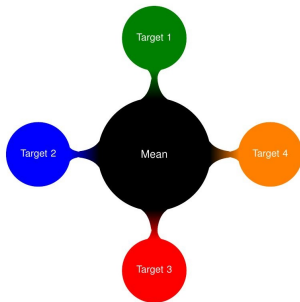


Advanced Machine Learning

Multi-Target Prediction: Methods Part 1



Learning goals

- Independent models for targets
- Mean regularization
- Stacking
- Weight sharing in DL

INDEPENDENT MODELS

- Assume a linear basis function model for the m -th target:

$$f_k(\mathbf{x}) = \boldsymbol{\theta}_k^\top \phi(\mathbf{x}),$$

θ_k is target-specific parameter and ϕ some feature mapping.

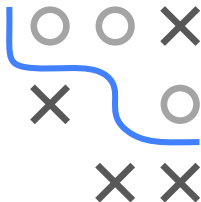
- Use this with with large nr of targets.
- We optimize jointly:

$$\min_{\Theta} \|Y - \Phi\Theta\|_F^2 + \sum_{m=1}^l \lambda_m \|\theta_m\|^2,$$

$$\|B\|_F^2 = \sqrt{\sum_{i=1}^n \sum_{m=1}^l B_{i,m}^2}$$
 is Frobenius norm for $B \in \mathbb{R}^{n \times l}$ and

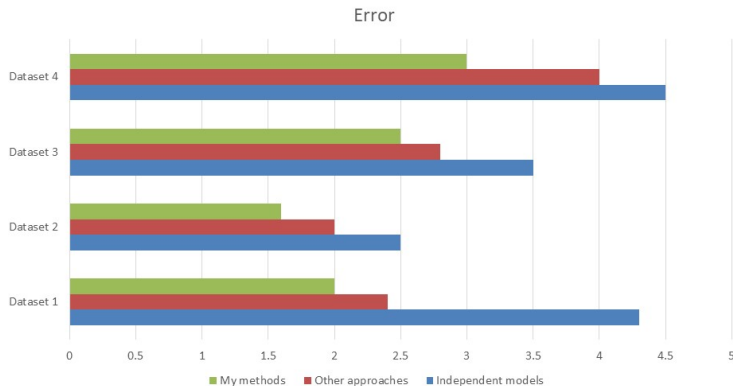
$$\Phi = \begin{bmatrix} \phi(\mathbf{x}^{(1)})^\top \\ \vdots \\ \phi(\mathbf{x}^{(n)})^\top \end{bmatrix} \quad \Theta = [\theta_1 \quad \cdots \quad \theta_l].$$

Frobenius norm = sum of SSE-s of all targets



INDEPENDENT MODELS

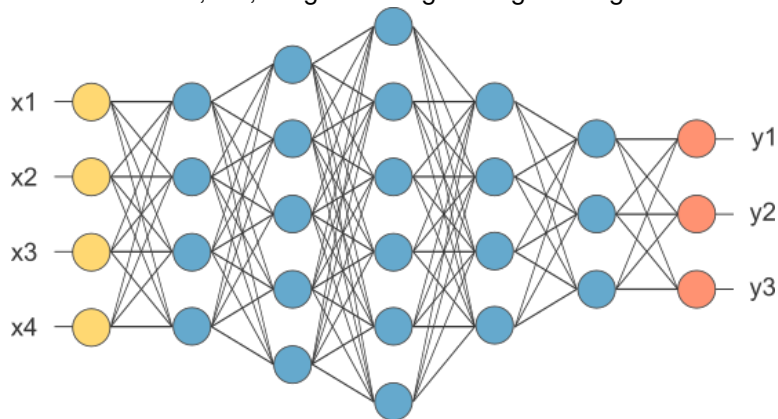
The experimental results section of a typical MTP paper:



→ Independent models don't exploit target deps, compared to more sophisticated methods, seems to be key for better performance.

ENFORCING SIMILARITY IN DEEP LEARNING

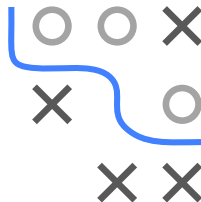
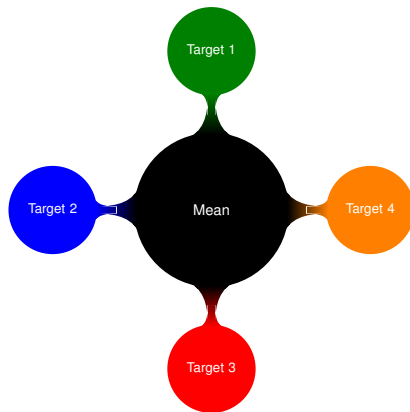
Commonly-used architecture: weight sharing in the final layer with m nodes, i.e., weight sharing among the targets



Caruana, Multitask learning: A knowledge-based source of inductive bias. Machine Learning 1997.

MEAN-REGULARIZED MULTI-TASK LEARNING

- Models for similar targets should behave similarly
- So params should be similar




- Approach: Bias parameter vectors towards mean vector:

$$\min_{\Theta} \|Y - \Phi\Theta\|_F^2 + \lambda \sum_{m=1}^l \|\theta_m - \frac{1}{l} \sum_{m'=1}^l \theta_{m'}\|^2$$

STACKING

- Originally, general ensemble learning technique.
- Level 1: apply series of ML methods on the same dataset
- Level 2: apply ML method to a new dataset consisting of the predictions obtained at level 1



Wolpert, Stacked generalization. Neural Networks 1992.

STACKING APPLIED TO MTP

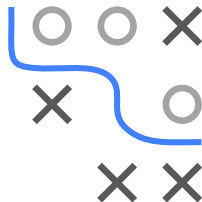
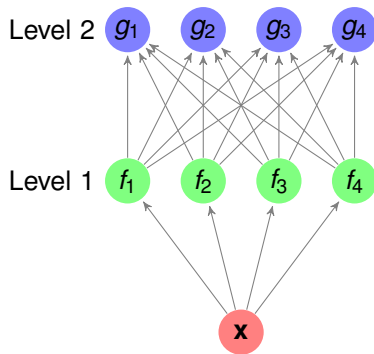
- Level 1: learn all $f_k(\mathbf{x})$ independently
- Level 2: learn model for each target independently, using predictions of level 1 \rightsquigarrow

$$f(\mathbf{x}) = g(f_1(\mathbf{x}), \dots, f_l(\mathbf{x}))$$

Or:

$$f(\mathbf{x}) = g(f_1(\mathbf{x}), \dots, f_l(\mathbf{x}), \mathbf{x})$$

- Advantages: easy to implement and general
- Has been shown to avoid overfitting in multivariate regression
- If level 2 learner uses regularization \rightsquigarrow models are forced to learn similar parameters for different targets.



Cheng and Hüllermeier, Combining Instance-based learning and Logistic Regression for Multi-Label classification, Machine

Learning, 2009.

STACKING VS BINARY RELEVANCE: EXAMPLE

- Compare F1-Score of random forest with stacking vs random forest with binary relevance on different multilabel datasets:

	birds	emotions	enron	genbase	image	langLog	reuters	scene	slashdot	yeast
BR(rf) F1-Score	0.637	0.620	0.578	0.989	0.431	0.319	0.671	0.616	0.441	0.615
STA(rf) F1-Score	0.646	0.634	0.583	0.986	0.446	0.317	0.685	0.633	0.453	0.624

- F1-Score is decomposed over targets.
- NB: Stacking slightly outperforms binary relevance on average.
- For more details, please refer to [Probst et al., 2017](#).

